

CS 501: AI and Information Retrieval

Lecture 1: Introduction to Retrieval-Augmented Generation - RAG

Explain how RAG works: RAG combines information retrieval with an LLM.

Documents are loaded, split into chunks, embedded, and stored in Chroma.

When a query is asked, relevant chunks are retrieved and passed to the LLM.

Lecture 2: Text Chunking and Text Splitters

What is a text splitter? A text splitter breaks large documents into small chunks.

`RecursiveCharacterTextSplitter` splits hierarchically by paragraphs, lines, and spaces. It uses `chunk_size` and `chunk_overlap` to maintain context across adjacent chunks.